

## MODULE-4

### TRANSACTIONS MANAGEMENT

What is transaction? Explain the ACID properties of transactions

#### **Transaction:**

A transaction can be defined as a group of tasks that form a single logical unit.

#### **For example:**

Suppose we want to withdraw Rs.100 from an account then we will follow following operations:

1. Check account balance
  2. If sufficient balance is present request for withdrawal.
  3. Get the money.
  4. Calculate  $\text{Balance} = \text{Balance} - 100$
  5. Update account with new balance
- The above mentioned four steps denote one transaction.
  - In a database, each transaction should maintain ACID property to meet the consistency and integrity of the database.

#### **ACID PROPERTIES:**

##### **1. Atomicity:**

- This property states that each transaction must be considered as a single unit and must be completed fully or not completed at all.
- No transaction in the database is left half completed.
- Database should be in a state either before the transaction execution or after the transaction execution. It should not be in the state of executing.

#### **For example:**

In the above mentioned withdrawal of money transaction all the five steps must be completed fully or none of the step is completed. Suppose if the transaction gets failed after step 3, then the customer will get the money but the balance will not be updated

accordingly. The state of database should be either at before ATM withdrawal (i.e. customer without withdrawn money) or after ATM withdrawal (i.e. customer with money and account updated). This will make the system in a consistent state.

## **2. Consistency:**

- The database must remain in consistent state after performing any transaction.
- No transaction should have any adverse effect on the data residing in the database.
- If the database was in a consistent state before the execution of a transaction, it must remain consistent after the execution of the transaction as well.

## **3. Durability:**


- The database should be durable enough to hold all its latest updates even if the system fails or restarts.
- If a transaction updates a chunk of data in a database and commits, then the database will hold the modified data.
- If a transaction commits but the system fails before the data could be written on to the disk, then that data will be updated once the system springs back into action.

## **4. Isolation**

- In a database system where more than one transaction are being executed simultaneously and in parallel, the property of isolation states that all the transactions will be carried out and executed as if it is the only transaction in the system.
- No transaction will affect the existence of any other transaction.

# TRANSACTION STATE

## States of Transaction:



### Active state

- The active state is the first state of every transaction. In this state, the transaction is being executed.
- For example: Insertion or deletion or updating a record is done here. But all the records are still not saved to the database.

### Partially committed

- In the partially committed state, a transaction executes its final operation, but the data is still not saved to the database.
- In the total mark calculation example, a final display of the total marks step is executed in this state.

### Committed

A transaction is said to be in a committed state if it executes all its operations successfully. In this state, all the effects are now permanently saved on the database system.

### **Failed state**

- If any of the checks made by the database recovery system fails, then the transaction is said to be in the failed state.
- In the example of total mark calculation, if the database is not able to fire a query to fetch the marks, then the transaction will fail to execute.

### **Aborted**

- If any of the checks fail and the transaction has reached a failed state then the database recovery system will make sure that the database is in its previous consistent state. If not then it will abort or roll back the transaction to bring the database into a consistent state.
- If the transaction fails in the middle of the transaction then before executing the transaction, all the executed transactions are rolled back to its consistent state.
- After aborting the transaction, the database recovery module will select one of the two operations:
  1. Re-start the transaction
  2. Kill the transaction

## **IMPLEMENTATION OF ATOMICITY AND DURABILITY USING SHADOW COPY**

- The recovery-management component of a database system can support atomicity and durability by a variety of schemes.
- Here we are going to learn about one of the simplest schemes called Shadow copy.

### **Shadow copy:**

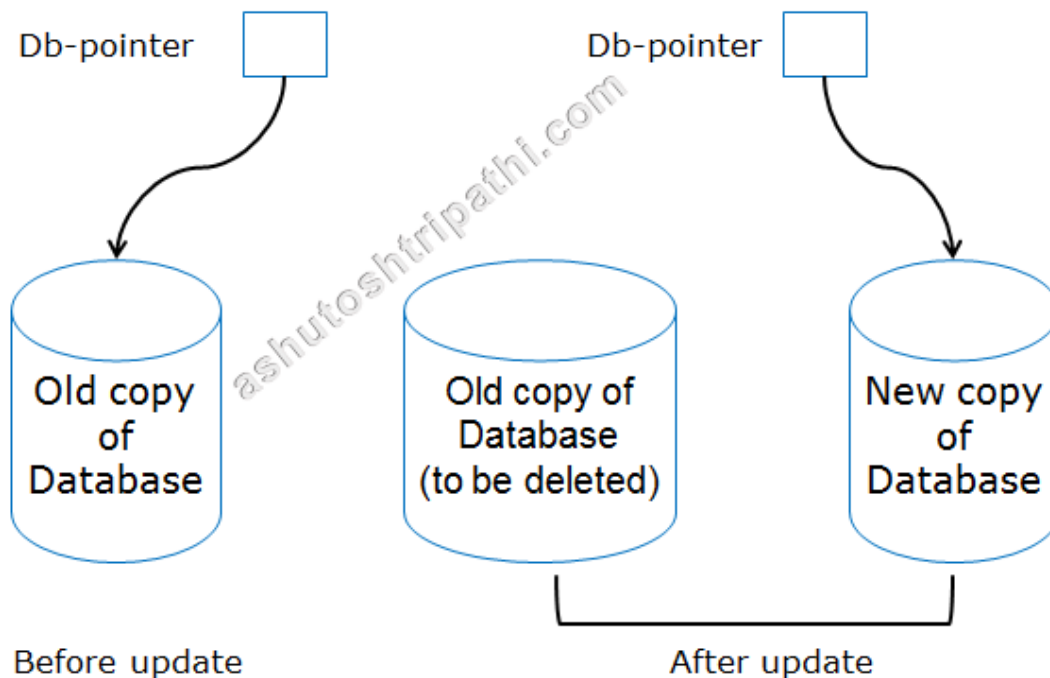
- In the shadow-copy scheme, a transaction that wants to update the database first creates a complete copy of the database. All updates are done on the new database copy, leaving the original copy, the shadow copy, untouched. If at any point the transaction has to be aborted, the system merely deletes the new copy. The old copy of the database has not been affected.

- This scheme is based on making copies of the database, called shadow copies, assumes that only one transaction is active at a time. The scheme also assumes that the database is simply a file on disk. A pointer called db-pointer is maintained on disk; it points to the current copy of the database.

If the transaction completes, it is committed as follows:

- First, the operating system is asked to make sure that all pages of the new copy of the database have been written out to disk. (Unix systems use the flush command for this purpose.)
- After the operating system has written all the pages to disk, the database system updates the pointer db-pointer to point to the new copy of the database; the new copy then becomes the current copy of the database. The old copy of the database is then deleted.

Figure below depicts the scheme, showing the database state before and after the update.



Shadow-copy technique for atomicity and durability

The transaction is said to have been committed at the point where the updated db pointer is written to disk.

#### **How the technique handles transaction failures:**

- If the transaction fails at any time before db-pointer is updated, the old contents of the database are not affected.
- We can abort the transaction by just deleting the new copy of the database.
- Once the transaction has been committed, all the updates that it performed are in the database pointed to by db pointer.
- Thus, either all updates of the transaction are reflected, or none of the effects are reflected, regardless of transaction failure.

#### **How the technique handles system failures:**

- Suppose that the system fails at any time before the updated db-pointer is written to disk. Then, when the system restarts, it will read db-pointer and will thus see the original contents of the database, and none of the effects of the transaction will be visible on the database.
- Next, suppose that the system fails after db-pointer has been updated on disk. Before the pointer is updated, all updated pages of the new copy of the database were written to disk.
- Again, we assume that, once a file is written to disk, its contents will not be damaged even if there is a system failure. Therefore, when the system restarts, it will read db-pointer and will thus see the contents of the database after all the updates performed by the transaction.

### **DBMS Concurrency Control**

Concurrency Control is the management procedure that is required for controlling concurrent execution of the operations that take place on a database.

But before knowing about concurrency control, we should know about concurrent execution.

## Concurrent Execution in DBMS

- In a multi-user system, multiple users can access and use the same database at one time, which is known as the concurrent execution of the database. It means that the same database is executed simultaneously on a multi-user system by different users.
- While working on the database transactions, there occurs the requirement of using the database by multiple users for performing different operations, and in that case, concurrent execution of the database is performed.
- The thing is that the simultaneous execution that is performed should be done in an interleaved manner, and no operation should affect the other executing operations, thus maintaining the consistency of the database. Thus, on making the concurrent execution of the transaction operations, there occur several challenging problems that need to be solved.

### Problems with Concurrent Execution

In a database transaction, the two main operations are **READ** and **WRITE** operations. So, there is a need to manage these two operations in the concurrent execution of the transactions as if these operations are not performed in an interleaved manner, and the data may become inconsistent. So, the following problems occur with the Concurrent Execution of the operations:

#### Problem 1: Lost Update Problems (W - W Conflict)

*The problem occurs* when two different database transactions perform the read/write operations on the same database items in an interleaved manner (i.e., concurrent execution) that makes the values of the items incorrect hence making the database inconsistent.

**For example:**

**Consider the below diagram where two transactions  $T_X$  and  $T_Y$ , are performed on the same account A where the balance of account A is \$300.**

Time	T <sub>x</sub>	T <sub>y</sub>
t <sub>1</sub>	READ (A)	—
t <sub>2</sub>	A = A - 50	
t <sub>3</sub>	—	READ (A)
t <sub>4</sub>	—	A = A + 100
t <sub>5</sub>	—	—
t <sub>6</sub>	WRITE (A)	—
t <sub>7</sub>		WRITE (A)

#### LOST UPDATE PROBLEM

- At time t<sub>1</sub>, transaction T<sub>x</sub> reads the value of account A, i.e., \$300 (only read).
- At time t<sub>2</sub>, transaction T<sub>x</sub> deducts \$50 from account A that becomes \$250 (only deducted and not updated/write).
- Alternately, at time t<sub>3</sub>, transaction T<sub>y</sub> reads the value of account A that will be \$300 only because T<sub>x</sub> didn't update the value yet.
- At time t<sub>4</sub>, transaction T<sub>y</sub> adds \$100 to account A that becomes \$400 (only added but not updated/write).
- At time t<sub>6</sub>, transaction T<sub>x</sub> writes the value of account A that will be updated as \$250 only, as T<sub>y</sub> didn't update the value yet.
- Similarly, at time t<sub>7</sub>, transaction T<sub>y</sub> writes the values of account A, so it will write as done at time t<sub>4</sub> that will be \$400. It means the value written by T<sub>x</sub> is lost, i.e., \$250 is lost.

Hence data becomes incorrect, and database sets to inconsistent.



### Dirty Read Problems (W-R Conflict)

The dirty read problem occurs *when one transaction updates an item of the database, and somehow the transaction fails, and before the data gets rollback, the updated database item is accessed by another transaction. There comes the Read-Write Conflict between both transactions.*

**For example:**

**Consider two transactions  $T_X$  and  $T_Y$  in the below diagram performing read/write operations on account A where the available balance in account A is \$300:**

Time	$T_X$	$T_Y$
$t_1$	READ (A)	—
$t_2$	$A = A + 50$	—
$t_3$	WRITE (A)	—
$t_4$	—	READ (A)
$t_5$	SERVER DOWN ROLLBACK	—

**DIRTY READ PROBLEM**

- At time  $t_1$ , transaction  $T_X$  reads the value of account A, i.e., \$300.
- At time  $t_2$ , transaction  $T_X$  adds \$50 to account A that becomes \$350.
- At time  $t_3$ , transaction  $T_X$  writes the updated value in account A, i.e., \$350.
- Then at time  $t_4$ , transaction  $T_Y$  reads account A that will be read as \$350.
- Then at time  $t_5$ , transaction  $T_X$  rolls back due to server problem, and the value changes back to \$300 (as initially).

- But the value for account A remains \$350 for transaction  $T_Y$  as committed, which is the dirty read and therefore known as the Dirty Read Problem.

### Unrepeatable Read Problem (W-R Conflict)

*Also known as Inconsistent Retrievals Problem that occurs when in a transaction, two different values are read for the same database item.*

**For example:**

**Consider two transactions,  $T_X$  and  $T_Y$ , performing the read/write operations on account A, having an available balance = \$300. The diagram is shown below:**

Time	$T_X$	$T_Y$
$t_1$	READ (A)	—
$t_2$	—	READ (A)
$t_3$	—	$A = A + 100$
$t_4$	—	WRITE (A)
$t_5$	READ (A)	—

**UNREPEATABLE READ PROBLEM**

- At time  $t_1$ , transaction  $T_X$  reads the value from account A, i.e., \$300.
- At time  $t_2$ , transaction  $T_Y$  reads the value from account A, i.e., \$300.
- At time  $t_3$ , transaction  $T_Y$  updates the value of account A by adding \$100 to the available balance, and then it becomes \$400.
- At time  $t_4$ , transaction  $T_Y$  writes the updated value, i.e., \$400.

- After that, at time  $t_5$ , transaction  $T_X$  reads the available value of account A, and that will be read as \$400.
- It means that within the same transaction  $T_X$ , it reads two different values of account A, i.e., \$ 300 initially, and after updating made by transaction  $T_Y$ , it reads \$400. It is an unrepeatable read and is therefore known as the Unrepeatable read problem.

Thus, in order to maintain consistency in the database and avoid such problems that take place in concurrent execution, management is needed, and that is where the concept of Concurrency Control comes into role.

## **Conflict Serializable Schedule**

- A schedule is called conflict serializability if after swapping of non-conflicting operations, it can transform into a serial schedule.
- The schedule will be a conflict serializable if it is conflict equivalent to a serial schedule.

## **Conflicting Operations**


The two operations become conflicting if all conditions satisfy:

1. Both belong to separate transactions.
2. They have the same data item.
3. They contain at least one write operation.

## **Example:**

Swapping is possible only if S1 and S2 are logically equal.

**1. T1: Read(A) T2: Read(A)**




**Schedule S1**

**Schedule S2**

Here,  $S1 = S2$ . That means it is non-conflict.

**2. T1: Read(A) T2: Write(A)**



**Schedule S1**

**Schedule S2**

Here,  $S1 \neq S2$ . That means it is conflict.

## Conflict Equivalent

In the conflict equivalent, one can be transformed to another by swapping non-conflicting operations. In the given example, S2 is conflict equivalent to S1 (S1 can be converted to S2 by swapping non-conflicting operations).

Two schedules are said to be conflict equivalent if and only if:

1. They contain the same set of the transaction.
2. If each pair of conflict operations are ordered in the same way.

**Example:**

### Non-serial schedule

T1	T2
Read(A) Write(A)	
	Read(A) Write(A)
Read(B) Write(B)	
	Read(B) Write(B)

### Schedule S1

### Serial Schedule

T1	T2
Read(A) Write(A) Read(B) Write(B)	
	Read(A) Write(A) Read(B) Write(B)

### Schedule S2

Schedule S2 is a serial schedule because, in this, all operations of T1 are performed before starting any operation of T2. Schedule S1 can be transformed into a serial schedule by swapping non-conflicting operations of S1.

**After swapping of non-conflict operations, the schedule S1 becomes:**

T1	T2
Read(A) Write(A) Read(B) Write(B)	     Read(A) Write(A) Read(B) Write(B)

Since, S1 is conflict serializable.

## **VIEW SERIALIZABILITY**

- A schedule will view serializable if it is view equivalent to a serial schedule.
- If a schedule is conflict serializable, then it will be view serializable.
- The view serializable which does not conflict serializable contains blind writes.

### **View Equivalent**

Two schedules S1 and S2 are said to be view equivalent if they satisfy the following conditions:

#### **1. Initial Read**

An initial read of both schedules must be the same. Suppose two schedule S1 and S2. In schedule S1, if a transaction T1 is reading the data item A, then in S2, transaction T1 should also read A.

T1	T2
Read(A)	Write(A)

**Schedule S1**

T1	T2
Read(A)	Write(A)

**Schedule S2**

Above two schedules are view equivalent because Initial read operation in S1 is done by T1 and in S2 it is also done by T1.

## 2. Updated Read

In schedule S1, if Ti is reading A which is updated by Tj then in S2 also, Ti should read A which is updated by Tj.

T1	T2	T3
Write(A)	Write(A)	Read(A)

**Schedule S1**

T1	T2	T3
Write(A)	Write(A)	<u>Read(A)</u>

**Schedule S2**

Above two schedules are not view equal because, in S1, T3 is reading A updated by T2 and in S2, T3 is reading A updated by T1.

### 3. Final Write

A final write must be the same between both the schedules. In schedule S1, if a transaction T1 updates A at last then in S2, final writes operations should also be done by T1.

T1	T2	T3
Write(A)	Read(A)	Write(A)

**Schedule S1**

T1	T2	T3
Write(A)	Read(A)	Write(A)

**Schedule S2**

Above two schedules is view equal because Final write operation in S1 is done by T3 and in S2, the final write operation is also done by T3.

#### Example:

T1	T2	T3
Read(A)	Write(A)	Write(A)
Write(A)		

**Schedule S**



With 3 transactions, the total number of possible schedule

1.  $= 3! = 6$
2.  $S1 = \langle T1\ T2\ T3 \rangle$
3.  $S2 = \langle T1\ T3\ T2 \rangle$
4.  $S3 = \langle T2\ T3\ T1 \rangle$
5.  $S4 = \langle T2\ T1\ T3 \rangle$
6.  $S5 = \langle T3\ T1\ T2 \rangle$
7.  $S6 = \langle T3\ T2\ T1 \rangle$

**Taking first schedule S1:**

<b>T1</b>	<b>T2</b>	<b>T3</b>
Read(A) Write(A)	Write(A)	Write(A)

**Schedule S1**

**Step 1:** final updating on data items

In both schedules S and S1, there is no read except the initial read that's why we don't need to check that condition.

**Step 2:** Initial Read

The initial read operation in S is done by T1 and in S1, it is also done by T1.

**Step 3:** Final Write

The final write operation in S is done by T3 and in S1, it is also done by T3. So, S and S1 are view Equivalent.

The first schedule S1 satisfies all three conditions, so we don't need to check another schedule.

**Hence, view equivalent serial schedule is:**

1. T1 → T2 → T3

## RECOVERABILITY OF SCHEDULE

Sometimes a transaction may not execute completely due to a software issue, system crash or hardware failure. In that case, the failed transaction has to be rollback. But some other transaction may also have used value produced by the failed transaction. So we also have to rollback those transactions.

T1	T1's buffer space	T2	T2's buffer space	Database
				A = 6500
Read(A);	A = 6500			A = 6500
A = A - 500;	A = 6000			A = 6500
Write(A);	A = 6000			A = 6000
		Read(A);	A = 6000	A = 6000
		A = A + 1000;	A = 7000	A = 6000
		Write(A);	A = 7000	A = 7000
		Commit;		
Failure Point				
Commit;				

The above table 1 shows a schedule which has two transactions. T1 reads and writes the value of A and that value is read and written by T2. T2 commits but later on, T1 fails. Due to the failure, we have to rollback T1. T2 should also be rollback because it reads the value written by T1, but

T2 can't be rollback because it already committed. So this type of schedule is known as irrecoverable schedule.

**Irrecoverable schedule:** The schedule will be irrecoverable if Tj reads the updated value of Ti and Tj committed before Ti commit.

<b>T1</b>	<b>T1's buffer space</b>	<b>T2</b>	<b>T2's buffer space</b>	<b>Database</b>
				A = 6500
Read(A);	A = 6500			A = 6500
A = A - 500;	A = 6000			A = 6500
Write(A);	A = 6000			A = 6000
		Read(A);	A = 6000	A = 6000
		A = A + 1000;	A = 7000	A = 6000
		Write(A);	A = 7000	A = 7000
Failure Point				
Commit;				
		Commit;		

The above table 2 shows a schedule with two transactions. Transaction T1 reads and writes A, and that value is read and written by transaction T2. But later on, T1 fails. Due to this, we have to rollback T1. T2 should be rollback because T2 has read the value written by T1. As it has not committed before T1 commits so we can rollback transaction T2 as well. So it is recoverable with cascade rollback.

### Recoverable with cascading rollback:

The schedule will be recoverable with cascading rollback if  $T_j$  reads the updated value of  $T_i$ . Commit of  $T_j$  is delayed till commit of  $T_i$ .

T1	T1's buffer space	T2	T2's buffer space	Database
				A = 6500
Read(A);	A = 6500			A = 6500
A = A - 500;	A = 6000			A = 6500
Write(A);	A = 6000			A = 6000
Commit;		Read(A);	A = 6000	A = 6000
		A = A + 1000;	A = 7000	A = 6000
		Write(A);	A = 7000	A = 7000
		Commit;		

The above Table 3 shows a schedule with two transactions. Transaction T1 reads and write A and commits, and that value is read and written by T2. So this is a cascade less recoverable schedule.

### Thomas write Rule

Thomas Write Rule provides the guarantee of serializability order for the protocol. It improves the Basic Timestamp Ordering Algorithm.

The basic Thomas write rules are as follows:

- If  $TS(T) < R\_TS(X)$  then transaction T is aborted and rolled back, and operation is rejected.
- If  $TS(T) < W\_TS(X)$  then don't execute the  $W\_item(X)$  operation of the transaction and continue processing.

- If neither condition 1 nor condition 2 occurs, then allowed to execute the WRITE operation by transaction  $T_i$  and set  $W\_TS(X)$  to  $TS(T)$ .

If we use the Thomas write rule then some serializable schedule can be permitted that does not conflict serializable as illustrate by the schedule in a given figure:

<b>T1</b>	<b>T2</b>
R(A)	W(A) Commit
W(A) Commit	

**Figure:** A Serializable Schedule that is not Conflict Serializable

In the above figure, T1's read and precedes T1's write of the same data item. This schedule does not conflict serializable.

Thomas write rule checks that T2's write is never seen by any transaction. If we delete the write operation in transaction T2, then conflict serializable schedule can be obtained which is shown in below figure.

<b>T1</b>	<b>T2</b>
R(A)	Commit
W(A) Commit	

**Figure:** A Conflict Serializable Schedule

